



Ranking Ultimate Frisbee Teams


Insight Bayesian Workshop

Shane Bussmann
February 20, 2018


About Me

- Senior Data Scientist at CiBO Technologies since December 2017
- Lead Data Scientist at Understory Weather, 2015 - 2017
- Health Data Science Fellow, Boston, Summer 2015
- Academic history
 - PhD Astronomy, 2010, University of Arizona
 - Postdoc at Harvard-Smithsonian CfA, 2010 - 2013
 - Postdoc at Cornell, 2013 - 2015

Astronomy




Understory



CiBO Technologies is
Planetary-Scale

Agricultural

Simulation Optimization

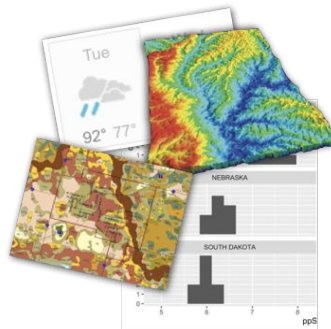


**Planetary Scale,
Daily Simulations**

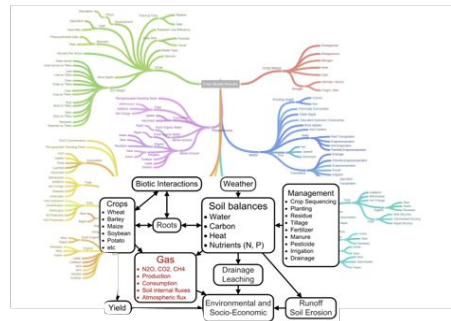
**Sub-field resolution
utilizing on-farm
and remote sensing**



**Multi-country,
Multi-crop**



**Thousands of
Variables
&
Millions of data points
*per field***



**Rich Domain
Model with Hundreds
of Interactions and
Outputs**



Goal of This Workshop

Walk through a real-world example to show how Bayesian analysis leads to better results!



The Problem

- BUDA (Boston Ultimate Disc Alliance) runs recreational ultimate frisbee leagues in the greater Boston area
- Teams self-assign to one of four divisions
 - Div 1 == highest
 - Div 3 == lowest
- Teams responsible for creating their own schedule
 - Some teams play lots of games, others not so much
 - Some teams play “out of division” games (e.g., Div 1 vs. Div 2)
- Some teams play with 4 men and 3 women, others with 5 men and 2 women
- We have been asked to rank teams for the end-of-season tournament



Ranking Teams: Plus/Minus?

- Sort by Plus/Minus
- No consideration of strength of schedule
- Is Upstream really better than AHOC?
- Who would win if they played each other?

Team Name	PlusMinus	divname	Wins	Losses	Ties	winpercent
Upstream	119	4/3 Div 2	20	6	0	0.769231
AHOC	100	4/3 Div 1	14	0	0	1.000000
Injustice League	100	4/3 Div 2	14	1	0	0.933333
Puddingstone	98	4/3 Div 2	18	6	0	0.750000
Crossroads	73	Open Div 1	17	3	1	0.850000
JuJu Hex	72	5/2 Div 2	11	3	1	0.785714
Flaming Croissants	66	4/3 Div 2	15	5	0	0.750000
Too Drunk to Fail	59	4/3 Div 3	11	2	0	0.846154
Pink Flamingos	58	5/2 Div 2	10	1	0	0.909091
Baba Yaga	47	5/2 Div 2	9	3	0	0.750000
Gothrilla	47	4/3 Div 1	12	4	1	0.750000
Jack's Abby HAOS Lager	46	4/3 Div 2	10	2	0	0.833333
License to Kilt (fka Scoobers in Scotland)	46	4/3 Div 2	13	6	0	0.684211



Ranking Teams: Division + Plus/Minus

- Sort by division
 - Then by Plus/Minus
- Worst Div 1 team > best Div 2 team
- Still ignoring strength of schedule
 - Are Upstream and Injustice League really worse than BBN and LPFK?

Team Name	PlusMinus	divname	Wins	Losses	Ties	winpercent
AHOC	100	4/3 Div 1	14	0	0	1.000000
Gothrilla	47	4/3 Div 1	12	4	1	0.750000
FlowChart	29	4/3 Div 1	12	4	0	0.750000
GrassBurner	29	4/3 Div 1	11	8	2	0.578947
Zerg Rush!	22	4/3 Div 1	10	4	0	0.714286
Swingers	-6	4/3 Div 1	5	7	1	0.416667
TuneSquad	-9	4/3 Div 1	6	7	1	0.461538
Turtle Boy	-12	4/3 Div 1	5	10	0	0.333333
Live Poultry, Fresh Killed (LPFK)	-39	4/3 Div 1	6	13	0	0.315789
BBN	-47	4/3 Div 1	2	8	1	0.200000
Upstream	119	4/3 Div 2	20	6	0	0.769231
Injustice League	100	4/3 Div 2	14	1	0	0.933333
Puddingstone	98	4/3 Div 2	18	6	0	0.750000
Flaming Croissants	66	4/3 Div 2	15	5	0	0.750000
Jack's Abby HAOS Lager	46	4/3 Div 2	10	2	0	0.833333
License to Kilt (fka Scoobers in Scotland)	46	4/3 Div 2	13	6	0	0.684211



Ranking Teams: A Bayesian Framework

- Suppose each team has a “rating” that represents the true skill level of that team
- Consider two teams, Team A and Team B, with ratings denoted as ratingA and ratingB , respectively
 - The difference in ratings is $\text{deltaRating} = \text{ratingA} - \text{ratingB}$
 - The probability that Team A wins is p
 - p is a function of deltaRating : $p = f(\text{deltaRating})$
- In our data, let’s assume that the outcome is binary (i.e., ignore ties) and is denoted “Team A Wins”
- Bayes rule tells us how to use our measurements (Team A Wins) to update our belief about deltaRating given our model for p

Team A	Team B	divname	Team A Wins
AHOC	Gothrilla	4/3 Div 1	True
AHOC	BBN	4/3 Div 1	True
AHOC	Stonecutters	4/3 Div 1	True
AHOC	FlowChart	4/3 Div 1	True
AHOC	Lady and the BAMF	4/3 Div 1	True
AHOC	Swingers	4/3 Div 1	True
AHOC	Upstream	4/3 Div 1	True
AHOC	Tubbs	4/3 Div 1	True
AHOC	Stonecutters	4/3 Div 1	True
AHOC	Zerg Rush!	4/3 Div 1	True
AHOC	Turtle Boy	4/3 Div 1	True
AHOC	Live Poultry, Fresh Killed (LPFK)	4/3 Div 1	True
AHOC	TuneSquad	4/3 Div 1	True
AHOC	JuJu Hex	4/3 Div 1	True
Gothrilla	AHOC	4/3 Div 1	False
Gothrilla	Swingers	4/3 Div 1	False



Ranking Teams: A Bayesian Framework

- Posterior (what we want)
 $P(f(\text{deltaRating}) \mid \text{Team A Wins})$
- Prior (where we start)
 $P(f(\text{deltaRating}))$
- Likelihood (how we go from the prior to the posterior)
 $P(\text{Team A Wins} \mid f(\text{deltaRating}))$
“What is the likelihood that team A won, given the rating differential between team A and B and our function that converts that differential to the probability that A wins”
- Support:
 $\text{Likelihood} / P(\text{Team A Wins}) = \text{Likelihood} / 0.5$

Team A	Team B	divname	Team A Wins
AHOC	Gothrilla	4/3 Div 1	True
AHOC	BBN	4/3 Div 1	True
AHOC	Stonecutters	4/3 Div 1	True
AHOC	FlowChart	4/3 Div 1	True
AHOC	Lady and the BAMF	4/3 Div 1	True
AHOC	Swingers	4/3 Div 1	True
AHOC	Upstream	4/3 Div 1	True
AHOC	Tubbs	4/3 Div 1	True
AHOC	Stonecutters	4/3 Div 1	True
AHOC	Zerg Rush!	4/3 Div 1	True
AHOC	Turtle Boy	4/3 Div 1	True
AHOC	Live Poultry, Fresh Killed (LPFK)	4/3 Div 1	True
AHOC	TuneSquad	4/3 Div 1	True
AHOC	JuJu Hex	4/3 Div 1	True
Gothrilla	AHOC	4/3 Div 1	False
Gothrilla	Swingers	4/3 Div 1	False



Ranking Teams: A Bayesian Framework

- Posterior (what we want)
 $P(f(\text{deltaRating}) \mid \text{Team A Wins})$
- Prior (where we start)
 $P(f(\text{deltaRating}))$
- Likelihood (how we go from the prior to the posterior)
 $P(\text{Team A Wins} \mid f(\text{deltaRating}))$
“What is the likelihood that team A won, given the rating differential between team A and B and our function that converts that differential to the probability that A wins”
- Support:
 $\text{Likelihood} / P(\text{Team A Wins}) = \text{Likelihood} / 0.5 \sim \text{Likelihood}$

$$P(f(\text{deltaRating}) \mid \text{Team A Wins}) = P(f(\text{deltaRating})) * P(\text{Team A Wins} \mid f(\text{deltaRating}))$$

Team A	Team B	divname	Team A Wins
AHOC	Gothrilla	4/3 Div 1	True
AHOC	BBN	4/3 Div 1	True
AHOC	Stonecutters	4/3 Div 1	True
AHOC	FlowChart	4/3 Div 1	True
AHOC	Lady and the BAMF	4/3 Div 1	True
AHOC	Swingers	4/3 Div 1	True
AHOC	Upstream	4/3 Div 1	True
AHOC	Tubbs	4/3 Div 1	True
AHOC	Stonecutters	4/3 Div 1	True
AHOC	Zerg Rush!	4/3 Div 1	True
AHOC	Turtle Boy	4/3 Div 1	True
AHOC	Live Poultry, Fresh Killed (LPFK)	4/3 Div 1	True
AHOC	TuneSquad	4/3 Div 1	True
AHOC	JuJu Hex	4/3 Div 1	True
Gothrilla	AHOC	4/3 Div 1	False
Gothrilla	Swingers	4/3 Div 1	False




$p = f(\text{deltaRating})$

- First, we need to define the function that converts rating differential between two teams into a probability that the first team wins
- Key requirements
 - If $\text{deltaRating} < 0$, $p = 0$ (i.e., Team B wins always)
 - If $\text{deltaRating} > 0$, $p = 1$ (i.e., Team A wins always)
 - If $\text{deltaRating} = 0$, $p = 0.5$
- Ideas?


$p = f(\text{deltaRating})$

- First, we need to define the function that converts rating differential between two teams into a probability that the first team wins
- Key requirements
 - If $\text{deltaRating} < 0$, $p = 0$ (i.e., Team B wins always)
 - If $\text{deltaRating} > 0$, $p = 1$ (i.e., Team A wins always)
 - If $\text{deltaRating} = 0$, $p = 0.5$
- sigmoid function
 - $y = 1 / (1 + \exp(-x))$
 - jupyter notebook: Insight Bayesian Workshop - Sigmoid



Priors

- To start out, let's assume we know nothing about each team's rating before any games have been played
- Every team is "average"
 - Normal distribution
 - Mean = 0
 - Standard deviation = 1



In PyMC3:

```
with pm.model() as model:  
    ratings = pm.Normal(  
        'ratings', mu=0, sd=1, shape=nTeams)
```

- The keyword `shape` means we will get a vector of ratings, one for each team in the dataframe.
- `nTeams` is the number of teams in the database



Building deltaRating

- We need `deltaRating` for every row in the dataframe (i.e., for every game in the season)
- Key insight: we can index `ratings` just like a normal python list
- Plan: assign a number to each team and use that number to index `ratings`

Team A	Team B	Team A Wins	Index A	Index B
AHOC	Gothrilla	True	0	1
AHOC	BBN	True	0	9
AHOC	Stonecutters	True	0	41
AHOC	FlowChart	True	0	2
AHOC	Lady and the BAMF	True	0	28

In PyMC3:

```
with pm.model() as model:
    ratings = pm.Normal(
        'ratings', mu=0, sd=1, shape=nTeams)

    deltaRating = ratings[indexA] -
                    ratings[indexB]
```



Implementing sigmoid

- $p = \text{sigmoid}(\text{deltaRating})$
- Recall that $\text{sigmoid} = 1 / (1 + \exp(-x))$

In PyMC3:

```
with pm.model() as model:
    ratings = pm.Normal(
        'ratings', mu=0, sd=1, shape=nTeams)

    deltaRating = ratings[indexA] -
                    ratings[indexB]

    p = 1 / (1 + np.exp(-deltaRating))
```



Likelihood Distribution

- $P(\text{Team A Wins} \mid f(\text{deltaRating}))$ means “What is the likelihood that team A won, given the rating differential between team A and B and our function that converts that differential to the probability that A wins”?
- Team A Wins is a binary outcome
- What distribution should we use to model our likelihood? What distribution links a probability to a binary outcome?

In PyMC3:

```
with pm.model() as model:
    ratings = pm.Normal(
        'ratings', mu=0, sd=1, shape=nTeams)

    deltaRating = ratings[indexA] -
                    ratings[indexB]

    p = 1 / (1 + np.exp(-deltaRating))
```




Likelihood Distribution

- $P(\text{Team A Wins} \mid f(\text{deltaRating}))$ means “What is the likelihood that team A won, given the rating differential between team A and B and our function that converts that differential to the probability that A wins”?
- Team A Wins is a binary outcome
- What distribution should we use to model our likelihood? What distribution links a probability to a binary outcome?
- Bernoulli Distribution:
<http://docs.pymc.io/api/distributions/discrete.html#pymc3.distributions.discrete.Bernoulli>

In PyMC3:

```
with pm.model() as model:
    ratings = pm.Normal(
        'ratings', mu=0, sd=1, shape=nTeams)

    deltaRating = ratings[indexA] -
                  ratings[indexB]

    p = 1 / (1 + np.exp(-deltaRating))

    win = pm.Bernoulli('win', p,
                      observed=team_A_wins)
```

That's it! We've specified our model. Now we can sample and get the posterior. Let's go to a notebook to see what happens!



Let's test it!

But first, we'll look at a Jupyter notebook for an example of a simulated season with only 3 teams of known rating.




What went wrong?

- Got the top 2 teams right, but many other instances of questionable results
- Ideas for changes?

What went wrong?

- Got the top 2 teams right, but many other instances of questionable results
- Ideas for changes?
 - Support?
 - Bernoulli distribution is definitely the right way to convert probability to binary outcome
 - Sigmoid may not be the best way to convert rating differential to probability, but it seems reasonable
 - Priors?
 - We expect that a Div 1 team should beat a Div 2 team. Maybe we should try to incorporate that prior knowledge into our model?
 - One possible approach:
 - Avg. Div 1 beats Avg. Div 2 about 95% of the time
 - Avg. Div 2 beats Avg. Div 3 about 75% of the time





Let's test it!

Switch to Jupyter notebook.



Conclusion

- We looked at different approaches to ranking teams in BUDA ultimate frisbee leagues
- Heuristic approaches are kind of sort of ok, but suffer from problems:
 - subjective
 - offer no clear path to validation
 - can't be extended in interesting ways
- In contrast, our Bayesian approach is quantitative, offers direct validation, and can be extended in really interesting ways (e.g., what is the likelihood of team A winning the tournament?)
- We looked at two different priors and showed that they have a significant impact on the results
 - This is ok because we stated at the outset very clearly what our priors were



Next steps

- Ideas
 - Simulate the end-of-season tournament
 - Use accuracy, logloss, or some other metric to optimize the model (watch out for overfitting!)
 - Use data from previous years:
 - to specify prior for this year
 - to implement cross-validation
 - Incorporate goals-for and goals-against into analysis (Binomial distribution)
- Game scores (for all seasons through 2016) and jupyter notebooks are available to you
- I'm very interested to see what you come up with!
 - sbussmann@cibotechnologies.com
- I will be at Boston Bayesians meetup on Feb 22, hope to see you there!



Resources

- <http://camdavidsonpilon.github.io/Probabilistic-Programming-and-Bayesian-Methods-for-Hackers/>
- PyMC3: <http://docs.pymc.io/#learn-bayesian-statistics-with-a-book-together-with-pymc3>
- PyMC3 discussion forum (very active, helpful userbase): <https://discourse.pymc.io/>

$p = f(\text{deltaRating}, \text{sharpness})$

- Next, we need to add a parameter that characterizes how sharp the transition is from $p=0$ to $p=0.5$ and from $p=0.5$ to $p=1.0$
- We will call this “sharpness”
- Higher sharpness => faster transition from $p=0$ to $p=0.5$ and from $p=0.5$ to $p=1.0$
- Lower sharpness => slower transition from $p=0$ to $p=0.5$ and from $p=0.5$ to $p=1.0$
- sigmoid function with sharpness
 - $y = (1 + \exp(-\text{sharpness} * x))$
 - jupyter notebook: Insight Bayesian Workshop - Sigmoid

